

Practical Lab - Computational Finance

Sheet 1

Min Liu, Christian Fiedler and Tim Griesbach

November 4th, 2019

Simulate the normal distribution

Task 1. *What does the given code do? What is the difference between the two ways to generate random numbers? What happens if you remove the expression `(double)` in the line marked with `//` *? Look at the GSL reference and see whether there is a direct function for simulating normally distributed random variables.*

The random number generator from GSL is a variant of the twisted generalized feedback shift-register algorithm (cited from reference). This random number generator pass also different statistical tests. `rand()` is a simpler functions. That means we could expect that the gsl random number generator causes less biased results than the method via `rand()`. The function `int rand(void)` returns a pseudo-random number in the range of 0 to `RAND_MAX`. `RAND_MAX` is a constant whose default value may vary between implementations but it is granted to be at least 32767.

Thus we can see `(double)rand()/RAND_MAX` as a random variable, which is unifomly distributed in $[0,1]$. Without "`(double)`" this number is just 0. This the case because the return type of `rand()` is `int`. So, this expression evaluates to zero because without the type casting the result has to be a `int`.

The rest codes set a uniform distributed random variable and free the memory by using `gsl_rng_free`. In the gsl reference we can find followings for the normal distribution:

- `double gsl_rng_gaussian(const gsl_rng * r, double sigma)` This function returns a Gaussian random variate, with mean zero and standard deviation `sigma`.
- `double gsl_rng_gaussian_pdf(double x, double sigma)` This function computes the probability density $p(x)$ at x for a Gaussian distribution with standard deviation `sigma`, using the formula given above.
- `double gsl_rng_ugaussian(const gsl_rng * r),`
`double gsl_rng_ugaussian_pdf(double x) and`
`double gsl_rng_ugaussian_ratio_method(const gsl_rng * r).` These functions

compute results for the unit Gaussian distribution. They are equivalent to the functions above with a standard deviation of one, $\sigma = 1$.

Task 2. *Implement the rejection sampling algorithm in C/C++ and output 1 000 000 values into a file.*

We draw uniformly distributed random variables in a interval $[a, b]$ using the function `gsl_rng_uniform` and a bijection between $[0, 1]$ and $[a, b]$.

The rejection sampling is just the implementation of the given steps with the interval $[-3, 3]$ in step 1 since $\int_{-3}^3 \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{x^2}{2}\right) \approx 0.9973$. A too large interval could increase the runtime because we could have more samples x' with big $|x'|$. That means $p(x')$ will be small and we need in the mean more loops pass until the rejection sampling returns a value.

Task 3. *Inspect your data using MATLAB. What does the MATLAB script do? Do you get a sensible result? What did probably go wrong in Fig. 1?*

The Matlab script reads the textfile `values.txt` in the current folder that should contain samples of a gaussian random variable.

These samples are assigned to the variable `data` and then by the `[vals, bins] = hist(data, 100)` the samples are split into hundred bins. `vals` and `bins` are vectors. `vals` contains the frequency counts and `bins` contains the locations of the bins (centers of the corresponding interval).

Then the standard gaussian density function is plotted and the frequency counts are also plotted. However, the frequency counts in `vals` are normalized by `trapz(bins, vals)` which approximates the integral of `vals` with respect to `bins`. The rest of the code just creates a legend and labels the axis.

Figure 1 shows the output of the given Matlab script. We can observe that the samples generated with our rejection sampling implementation fits good to the gaussian density. In figure 1 from the sheet probably the interval was chosen too small. This hypothesis is supported by the following plot in figure 2 with our implementation of the rejection sampling and the interval $[-2, 2]$.

Task 4. *How do you get a normally distributed random variable when you only have a uniform one? Not entirely by coincidence, the c.d.f. and inverse c.d.f. of the standard normal distribution are computed by the two algorithms given below. Implement both algorithms and write a program that draws standard normal distributed values.*

Claim 1. *For a random variable X , its cumulative distribution function $F_X : \mathbb{R} \rightarrow [0, 1]$ and the generalized inverse cumulative distribution function $F_X^{-1} : [0, 1] \rightarrow \mathbb{R} \cup \{-\infty, +\infty\}$ with $F_X^{-1}(y) = \inf_{x \in \mathbb{R}} \{F_X(x) \geq y\}$ holds: If $U \sim \text{Unif}([0, 1])$, then*

$$F_X^{-1}(U) \sim X. \quad (1)$$

Proof. We define $\tilde{X} := F_X^{-1}(U)$. For $x \in \mathbb{R}$ we get

$$F_{\tilde{X}}(x) = P[F_X^{-1}(U) \leq x] = P\left[\inf_{x' \in \mathbb{R}} \{F_X(x') \geq U\} \leq x\right] \quad (2)$$

$$= P[U \leq F_X(x)] = F_X(x). \quad (3)$$




Figure 1: Output of the given Matlab script from task 3 if we use 1 000 000 samples in our rejection sampling implementation.

□

The implementation of the approximation algorithms for the c.d.f. and the inverse c.d.f. were already given on the sheet.

We can just apply the inverse c.d.f to draw standard normal distributed random variables to a realization of a standard uniform distributed random variable. See figure 3 for a result.

Task 5. *Explain the general idea behind Moro's algorithm!*

According to the sheet Moro's algorithm is an approximation to the c.d.f. of the standard normal distribution.

The general idea behind the given algorithm is to approximate the c.d.f. of the standard normal piecewise with a rational functions that are evaluated by Horner's method (linear number of multiplications and summations in polynomials degrees).

For negative evaluation points x the algorithm uses that for $X \sim N(0, 1)$

$$F_X(x) = 1 - F_X(-x) \quad (4)$$

holds. Now for different remaining ranges are different rational functions used to approximate the c.d.f..

Task 6. *Implement the Box-Muller method and plot 1 000 samples in a 2d-plot.*

For the implementation we use the given expressions for Z_1, Z_2 from the sheet and generate standard uniform samples with the GSL library.

In figure 4 we can observe that the samples are concentrated in a circle with radius two around $(0, 0)$. That is exactly what we expect if we look at the density function of (Z_1, Z_2) . This observation become clearer if we consider the 10 000 samples with the same seed as in figure 5.




Figure 2: Output of the given Matlab script from task 3 if we use 1 000 000 samples in our rejection sampling implementation and the interval $[-2, 2]$.

Task 7. Do some research and be prepared to explain why z_1 and z_2 are standard normally distributed.

Claim 2. If $U_1, U_2 \sim \text{Unif}([0, 1])$ are independent, then $Z_1, Z_2 \sim N(0, 1)$ i.i.d. where

$$Z_1 := \sqrt{-2 \log U_1} \cos(2\pi U_2), \quad (5)$$

$$Z_2 := \sqrt{-2 \log U_1} \sin(2\pi U_2). \quad (6)$$

Proof. Let U_1 and U_2 be random variables with $U_1, U_2 \sim \text{Unif}([0, 1])$. We define

$$R := \sqrt{-2 \log U_1} \quad \text{and} \quad \Phi := 2\pi U_2. \quad (7)$$

Then we apply change of variable to the transformation

$$\psi : \mathbb{R}_{>0} \times [0, 2\pi) \rightarrow \mathbb{R} \times \mathbb{R}, \quad (r, \phi) \mapsto (r \cos \phi, r \sin \phi) \quad (8)$$

and we obtain that for the density function of

$$(Z_1, Z_2) := (R \cos \Phi, R \sin \Phi) \quad (9)$$

holds for $x = r \cos \phi$ and $y = r \sin \phi$

$$f_{Z_1, Z_2}(x, y) = \underbrace{f_{R, \Phi}(r, \phi)}_{\stackrel{(*)}{=} \frac{1}{2\pi} r \exp\left(-\frac{r^2}{2}\right)} \underbrace{\frac{1}{\det D\psi}}_{=\frac{1}{r}} = \frac{1}{2\pi} \exp\left(-\frac{x^2 + y^2}{2}\right) \quad (10)$$

The equality (\star) is satisfied because for $r \in \mathbb{R}_{>0}$

$$P[R \leq r] = P[\sqrt{-2 \log U_1} \leq r] = P\left[U_1 \geq \exp\left(-\frac{r^2}{2}\right)\right] = 1 - \exp\left(-\frac{r^2}{2}\right). \quad (11)$$




Figure 3: Output of the given Matlab script from task 3 if we use 1 000 000 samples in our simulation of a standard normal random variable using the idea explained in our answer to task 4.

That is why we get as density function of the random variable R

$$f_R(r) = r \exp\left(-\frac{r^2}{2}\right). \quad (12)$$

By assumptions U_1 and U_2 are independent that implies that R and Φ are independent. That means with $\Phi \sim \text{Unif}([0, 2\pi))$ we obtain for $r \in \mathbb{R}_{>0}$ and $\phi \in [0, 2\pi)$

$$f_{R,\Phi}(r, \phi) = r \exp\left(-\frac{r^2}{2}\right) \frac{1}{2\pi} \quad (13)$$

as density function of (R, Φ) .

By factorization of the density function in equation 10 we get that Z_1 and Z_2 are independent and $Z_1, Z_2 \sim N(0, 1)$. \square

Parameter estimation

Task 8. Compute 15 naively and show experimentally that it yields the same result as the given algorithm. What is the advantage of this algorithm?

$$\hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i \quad (14)$$

$$\hat{\sigma}^2 = \frac{1}{N-1} \sum_{i=1}^N (x_i - \hat{\mu})^2 \quad (15)$$

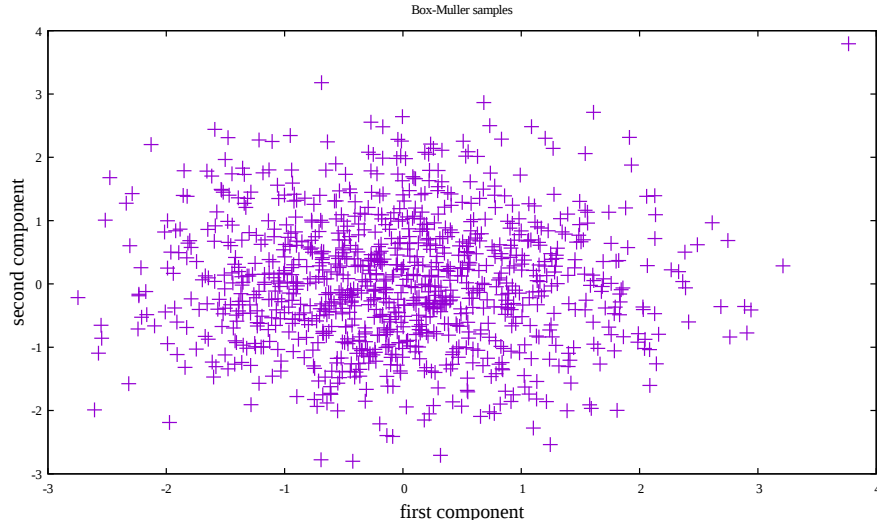


Figure 4: 2d plot of the 1 000 samples generated with the box muller method.

They can get the same results, since α in the i -th for-loop is just the mean value of the first i components and sigma is just the variation of first i components.

The second is better, since it decompose more complex computation into several easy computation such that rounding errors are less problematic.

Task 9. Choose an arbitrary μ and three different values for σ . Simulate N (up to $N_{\max} = 10\,000\,000$) samples by (6) and calculate $\hat{\sigma}$ for each of the three values of σ . Now generate a loglog- convergence plot (e.g. with GNU PLOT, MATLAB, matplotlib, pgfplots, . . .) of the estimation error $|\sigma - \hat{\sigma}|$ (for each σ) with respect to the number of samples N used for parameter estimation in the fashion of Fig. 2. How do you interpret the result?

In our convergence plot we see that for bigger σ the estimation error is bigger (with a few exceptions). This coincides roughly with the observations from figure 1 and 2 from the sheet. Furthermore, we can observe a linear decreasing trend in the plot and since figure 6 has two log-scaled axis the estimation error decreases exponentially in the number of samples N (see also the plotted function $l(N) = N^{-0.6}$).

Simulating and calibrating geometric Brownian motions

Task 10. We set $S(0) = 10$, $\mu = 0.1$, $\sigma = 0.2$, $T = 2$. Simulate 3 paths of a Wiener process with (2) and the corresponding asset prices (3) for $\Delta t = 0.5$ and 3 paths for $\delta t = 0.01$. Do the paths in principal look the same, independent of the time discretization? Is the mean change of the price $S(t)$ approximately $+0.1/\text{year}$?

The paths do not look in the same (compare figures 7 and 8), independent of the time steps. This is the case because we only get a very rough approximation (piecewise) of one




Figure 5: 2d plot of the 10 000 samples generated with the box muller method. The seed is the same as in figure 4.

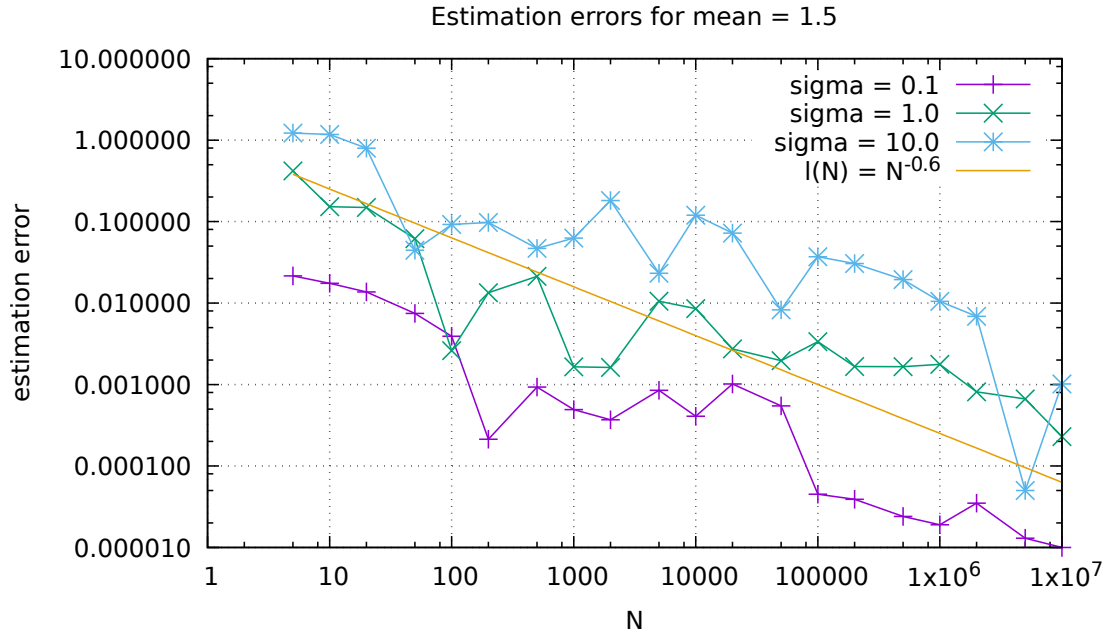


Figure 6: A plot of the estimation error with different σ and a fixed μ . The axis are log-scaled. Moreover, N is the number of samples.

sample and the number evaluations of the single standard normal distributed random variables differs. That is we do not get the same value for the time 1 even with both

time steps we get a value for this time.

The mean change of the price $S(t)$ holds roughly in figure 8. We can see this because of the plottes mean ($10 \exp(0.1x)$). However, this holds only roughly. This can be explained by the fact that for two years and a mean of 0.1 is a standard deviation of 0.2 quite high. Furthermore three paths are not enough to make a profound statement.

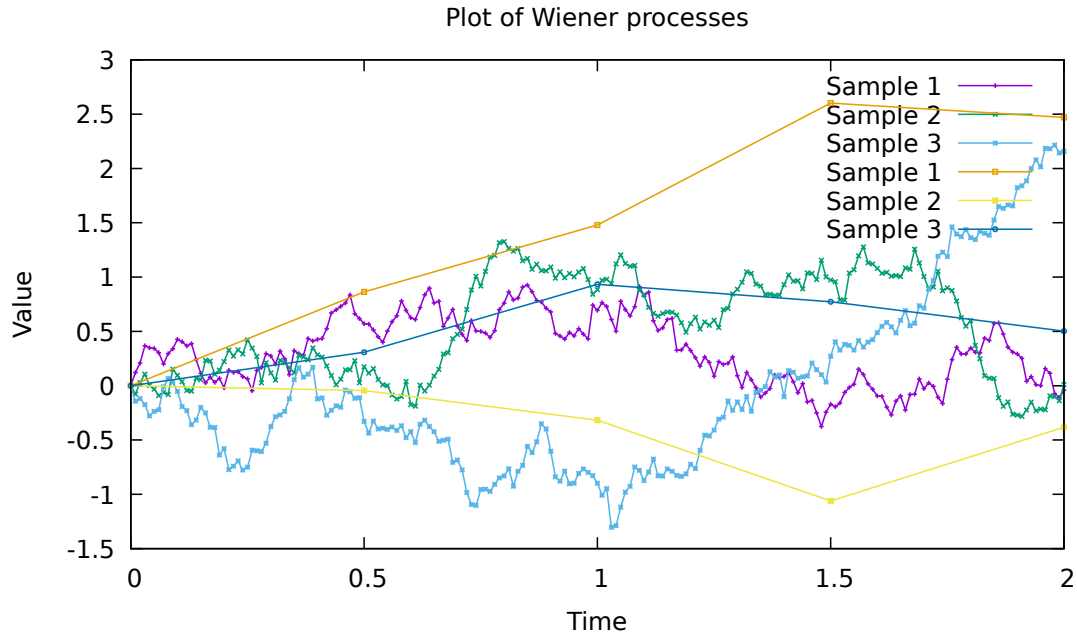


Figure 7: A plot three paths of the Wiener process with $\mu = 0.1$ and $\sigma = 0.2$ for different time steps (0.5 and 0.01).

Task 11. *Similar to Task 10, simulate a geometric Brownian motion with parameters $S(0) = 10$, $\mu = 0.1$, $\sigma = 0.2$, $T = 1$ and $\Delta t = 10^{-3}$. Given the values of $S(0)$, T and Δt , estimate the values of μ and σ from the 1000 time steps of your simulation. Explain your approach.*

The idea behind our approach is that the log-returns are normally distributed. In particular,

$$\log \left(\frac{S(t)}{S(0)} \right) = \left(\mu - \frac{1}{2} \sigma^2 \right) t + \sigma W(t). \quad (16)$$

This means $\log \left(\frac{S(t)}{S(0)} \right) \sim N \left(\left(\mu - \frac{1}{2} \sigma^2 \right) t, \sigma^2 t \right)$. Now we want to get rid of the dependence from t . That is why we consider

$$\log \left(\frac{S(t_i)}{S(t_{i-1})} \right) \sim N \left(\left(\mu - \frac{1}{2} \sigma^2 \right) (t_i - t_{i-1}), \sigma^2 (t_i - t_{i-1}) \right). \quad (17)$$

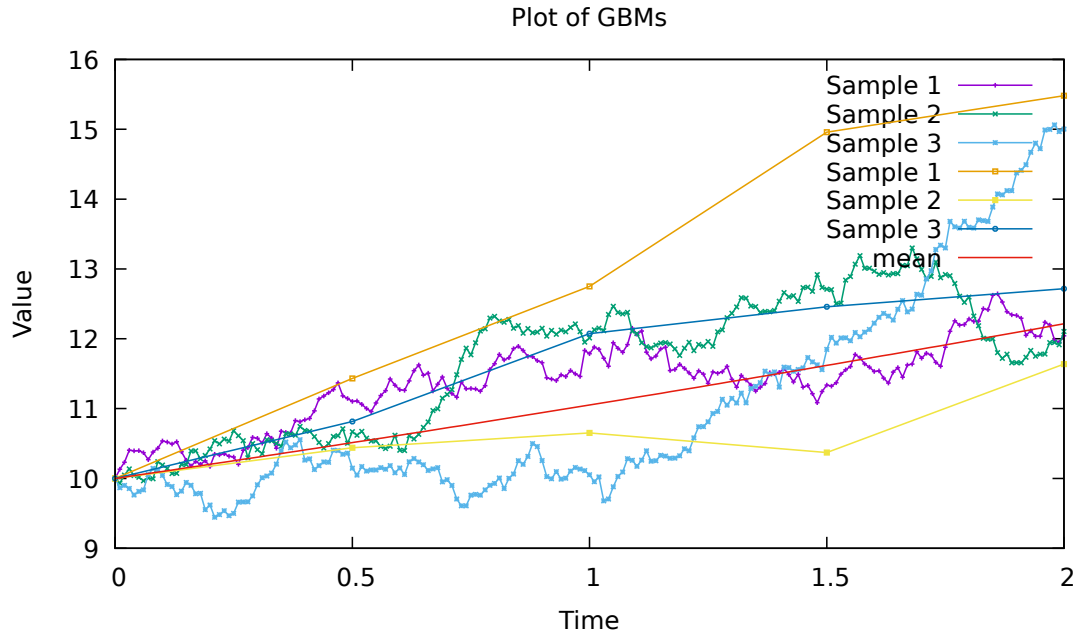


Figure 8: A plot three paths of the geometric Brownian motion with $S(0) = 10$, $\mu = 0.1$ and $\sigma = 0.2$ for different time steps (0.5 and 0.01).

Now estimate μ and σ as in task 8 what gives estimates for σ and μ after rewrite the equality of the log-return estimated and the above derived mean and standard deviation for time steps.

This approach is also implemented in our code.